

# Doing things with data frames

# Doing things with data frames

Let's go back to our Australian athletes:

```
##  
## -- Column specification -----  
## cols(  
##   Sex = col_character(),  
##   Sport = col_character(),  
##   RCC = col_double(),  
##   WCC = col_double(),  
##   Hc = col_double(),  
##   Hg = col_double(),  
##   Ferr = col_double(),  
##   BMI = col_double(),  
##   SSF = col_double(),  
##   `Bfat` = col_double(),  
##   LBM = col_double(),  
##   Ht = col_double()
```

# Choosing a column

```
athletes %>% select(Sport)
```

```
## # A tibble: 202 x 1
##   Sport
##   <chr>
## 1 Netball
## 2 Netball
## 3 Netball
## 4 Netball
## 5 Netball
## 6 Netball
## 7 Netball
## 8 Netball
## 9 Netball
## 10 Netball
## # ... with 192 more rows
```

# Choosing several columns

```
athletes %>% select(Sport, Hg, BMI)
```

```
## # A tibble: 202 x 3
##   Sport      Hg    BMI
##   <chr>    <dbl> <dbl>
## 1 Netball  13.6  19.2
## 2 Netball  12.7  21.2
## 3 Netball  12.3  21.4
## 4 Netball  12.3  21.0
## 5 Netball  12.8  21.8
## 6 Netball  11.8  21.4
## 7 Netball  12.7  21.5
## 8 Netball  12.4  24.4
## 9 Netball  12.4  22.6
## 10 Netball 14.1  22.8
## # ... with 192 more rows
```

# Choosing consecutive columns

```
athletes %>% select(Sex:WCC)
```

```
## # A tibble: 202 x 4
##   Sex      Sport      RCC      WCC
##   <chr>   <chr>   <dbl> <dbl>
## 1 female Netball  4.56  13.3
## 2 female Netball  4.15    6
## 3 female Netball  4.16   7.6
## 4 female Netball  4.32   6.4
## 5 female Netball  4.06   5.8
## 6 female Netball  4.12   6.1
## 7 female Netball  4.17    5
## 8 female Netball  3.8    6.6
## 9 female Netball  3.96   5.5
## 10 female Netball 4.44   9.7
## # ... with 192 more rows
```

## Choosing all-but some columns

```
athletes %>% select(-(RCC:LBM))
```

```
## # A tibble: 202 x 4
##   Sex      Sport      Ht      Wt
##   <chr>   <chr>   <dbl> <dbl>
## 1 female Netball  177.   59.9
## 2 female Netball  173.    63
## 3 female Netball  176    66.3
## 4 female Netball  170.   60.7
## 5 female Netball  183    72.9
## 6 female Netball  178.   67.9
## 7 female Netball  177.   67.5
## 8 female Netball  174.   74.1
## 9 female Netball  174.   68.2
## 10 female Netball 174.   68.8
## # ... with 192 more rows
```

# Select-helpers

Other ways to select columns: those whose name:

- `starts_with` something
- `ends_with` something
- `contains` something
- matches a “regular expression”
- `everything()` select all the columns

## Columns whose names begin with S

```
athletes %>% select(starts_with("S"))
```

```
## # A tibble: 202 x 3
##   Sex      Sport      SSF
##   <chr>   <chr>   <dbl>
## 1 female Netball    49
## 2 female Netball  110.
## 3 female Netball    89
## 4 female Netball  98.3
## 5 female Netball  122.
## 6 female Netball  90.4
## 7 female Netball  107.
## 8 female Netball  157.
## 9 female Netball  101.
## 10 female Netball  126.
## # ... with 192 more rows
```



# Columns whose names end with C

either uppercase or lowercase:

```
athletes %>% select(ends_with("c"))
```

```
## # A tibble: 202 x 3
##       RCC      WCC      Hc
##   <dbl> <dbl> <dbl>
## 1  4.56  13.3  42.2
## 2  4.15    6   38
## 3  4.16   7.6  37.5
## 4  4.32   6.4  37.7
## 5  4.06   5.8  38.7
## 6  4.12   6.1  36.6
## 7  4.17    5   37.4
## 8  3.8    6.6  36.5
## 9  3.96   5.5  36.3
## 10 4.44   9.7  41.4
```

# Case-sensitive

This works with any of the select-helpers:

```
athletes %>% select(ends_with("C", ignore.case=F))
```

```
## # A tibble: 202 x 2
```

```
##       RCC      WCC
```

```
##    <dbl> <dbl>
```

```
##  1  4.56  13.3
```

```
##  2  4.15    6
```

```
##  3  4.16   7.6
```

```
##  4  4.32   6.4
```

```
##  5  4.06   5.8
```

```
##  6  4.12   6.1
```

```
##  7  4.17    5
```

```
##  8  3.8    6.6
```

```
##  9  3.96   5.5
```

```
## 10  4.44   9.7
```

# Column names containing letter R

```
athletes %>% select(contains("r"))
```

```
## # A tibble: 202 x 3
##   Sport      RCC  Ferr
##   <chr>    <dbl> <dbl>
## 1 Netball  4.56    20
## 2 Netball  4.15    59
## 3 Netball  4.16    22
## 4 Netball  4.32    30
## 5 Netball  4.06    78
## 6 Netball  4.12    21
## 7 Netball  4.17   109
## 8 Netball  3.8     102
## 9 Netball  3.96    71
## 10 Netball 4.44    64
## # ... with 192 more rows
```

# Exactly two characters, ending with T

In regular expression terms, this is `^.t$`:

- `^` means “start of text”
- `.` means “exactly one character, but could be anything”
- `$` means “end of text”.

```
athletes %>% select(matches("^.t$"))
```

```
## # A tibble: 202 x 2
```

```
##       Ht      Wt
```

```
##    <dbl> <dbl>
```

```
##  1  177.   59.9
```

```
##  2  173.    63
```

```
##  3  176   66.3
```

```
##  4  170.   60.7
```

```
##  5  183   72.9
```

```
##  6  178.   67.9
```

```
##  7  177   67.5
```

# Choosing columns by property

- Use where as with summarizing several columns
- eg, to choose text columns:

```
athletes %>% select(where(is.character))
```

```
## # A tibble: 202 x 2
##   Sex      Sport
##   <chr>   <chr>
## 1 female Netball
## 2 female Netball
## 3 female Netball
## 4 female Netball
## 5 female Netball
## 6 female Netball
## 7 female Netball
## 8 female Netball
## 9 female Netball
```

# Choosing rows by number

```
athletes %>% slice(16:25)
```

```
## # A tibble: 10 x 13
```

```
##   Sex      Sport    RCC    WCC    Hc    Hg  Ferr  BMI
##   <chr>   <chr>   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 female Netba~  4.25  10.7  39.5  13.2   127  24.5
## 2 female Netba~  4.46  10.9  39.7  13.7   102  24.0
## 3 female Netba~  4.4    9.3  40.4  13.6    86  26.2
## 4 female Netba~  4.83   8.4  41.8  13.4    40  20.0
## 5 female Netba~  4.23   6.9  38.3  12.6    50  25.7
## 6 female Netba~  4.24   8.4  37.6  12.5    58  25.6
## 7 female Netba~  3.95   6.6  38.4  12.8    33  19.9
## 8 female Netba~  4.03   8.5  37.7  13     51  23.4
## 9 female BBall   3.96   7.5  37.5  12.3    60  20.6
## 10 female BBall   4.41   8.3  38.2  12.7    68  20.7
## # ... with 5 more variables: SSF <dbl>,
```

```
## #   %Defat <dbl>   TDM <dbl>   U+ <dbl>   U+ <dbl>
```

## Non-consecutive rows

```
athletes %>%  
  slice(10,13,17,42)
```

```
## # A tibble: 4 x 13  
##   Sex      Sport      RCC      WCC      Hc      Hg      Ferr      BMI  
##   <chr>   <chr>    <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>  
## 1 female Netball  4.44   9.7  41.4  14.1    64  22.8  
## 2 female Netball  4.02   9.1  37.7  12.7   107  23.0  
## 3 female Netball  4.46  10.9  39.7  13.7   102  24.0  
## 4 female Row      4.37   8.1  41.8  14.3    53  23.5  
## # ... with 5 more variables: SSF <dbl>,  
## #   %Bfat <dbl>, LBM <dbl>, Ht <dbl>, Wt <dbl>
```

## A random sample of rows

```
athletes %>% slice_sample(n=8)
```

```
## # A tibble: 8 x 13
##   Sex      Sport      RCC      WCC      Hc      Hg      Ferr      BMI
##   <chr>   <chr>   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 male    TSprnt    5.59   7.9   49.7   17.2   220   23.6
## 2 male    Row       5.04   7.1   44     14.8    64   25.8
## 3 male    Row       5.09  10.1  44.9   14.8   118   26.8
## 4 female T400m     4.27   6.9   44.1   14.7    45   20.4
## 5 female Netball  4.24   8.4   37.6   12.5    58   25.6
## 6 male    Swim      5.33   5.2   47.8   16.1   176   21.4
## 7 female Netball  4.46  10.9   39.7   13.7   102   24.0
## 8 male    BBall     4.73   6.7   42.8   14.9     8   19.8
## # ... with 5 more variables: SSF <dbl>,
## #   %Bfat <dbl>, LBM <dbl>, Ht <dbl>, Wt <dbl>
```



# Rows for which something is true

```
athletes %>% filter(Sport == "Tennis")
```

```
## # A tibble: 11 x 13
##   Sex      Sport    RCC    WCC    Hc    Hg  Ferr  BMI
##   <chr> <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 female Tennis    4    4.2  36.6  12    57  25.4
## 2 female Tennis  4.4    4   40.8  13.9   73  22.1
## 3 female Tennis  4.38   7.9  39.8  13.5   88  21.2
## 4 female Tennis  4.08   6.6  37.8  12.1  182  20.5
## 5 female Tennis  4.98   6.4  44.8  14.8   80  17.1
## 6 female Tennis  5.16   7.2  44.3  14.5   88  18.3
## 7 female Tennis  4.66   6.4  40.9  13.9  109  18.4
## 8 male   Tennis  5.66   8.3  50.2  17.7   38  23.8
## 9 male   Tennis  5.03   6.4  42.7  14.3  122  22.0
## 10 male   Tennis  4.97   8.8  43    14.9  233  22.3
## 11 male   Tennis  5.38   6.3  46    15.7   32  21.1
## # ... with 5 more variables: SSF <dbl>,
## #   %Bfat <dbl>, LBM <dbl>, Ht <dbl>, Wt <dbl>
```

## More complicated selections

```
athletes %>% filter(Sport == "Tennis", RCC < 5)
```

```
## # A tibble: 7 x 13
```

```
##   Sex      Sport    RCC    WCC    Hc    Hg  Ferr  BMI
##   <chr> <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 female Tennis    4     4.2  36.6  12     57  25.4
## 2 female Tennis  4.4     4   40.8  13.9   73  22.1
## 3 female Tennis  4.38   7.9  39.8  13.5   88  21.2
## 4 female Tennis  4.08   6.6  37.8  12.1  182  20.5
## 5 female Tennis  4.98   6.4  44.8  14.8   80  17.1
## 6 female Tennis  4.66   6.4  40.9  13.9  109  18.4
## 7 male    Tennis  4.97   8.8  43    14.9  233  22.3
## # ... with 5 more variables: SSF <dbl>,
## #   %Bfat <dbl>, LBM <dbl>, Ht <dbl>, Wt <dbl>
```

## Another way to do “and”

```
athletes %>% filter(Sport == "Tennis") %>%  
  filter(RCC < 5)
```

```
## # A tibble: 7 x 13
```

```
##   Sex      Sport    RCC    WCC    Hc    Hg    Ferr    BMI  
##   <chr> <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>  
## 1 female Tennis    4      4.2  36.6  12      57  25.4  
## 2 female Tennis  4.4      4  40.8  13.9    73  22.1  
## 3 female Tennis  4.38    7.9  39.8  13.5    88  21.2  
## 4 female Tennis  4.08    6.6  37.8  12.1   182  20.5  
## 5 female Tennis  4.98    6.4  44.8  14.8    80  17.1  
## 6 female Tennis  4.66    6.4  40.9  13.9   109  18.4  
## 7 male    Tennis  4.97    8.8  43    14.9   233  22.3  
## # ... with 5 more variables: SSF <dbl>,  
## #   %Bfat <dbl>, LBM <dbl>, Ht <dbl>, Wt <dbl>
```

## Either/Or

```
athletes %>% filter(Sport == "Tennis" | RCC > 5)
```

```
## # A tibble: 66 x 13
```

[illegible]







# “The top ones”

```
athletes %>%  
  arrange(desc(Wt)) %>%  
  slice(1:7) %>%  
  select(Sport, Wt)
```

```
## # A tibble: 7 x 2  
##   Sport      Wt  
##   <chr> <dbl>  
## 1 Field   123.  
## 2 BBall   114.  
## 3 Field   111.  
## 4 Field   108.  
## 5 Field   103.  
## 6 WPolo    101  
## 7 BBall    100.
```



## Another way

```
athletes %>%  
  slice_max(order_by = Wt, n=7) %>%  
  select(Sport, Wt)
```

```
## # A tibble: 7 x 2  
##   Sport      Wt  
##   <chr> <dbl>  
## 1 Field   123.  
## 2 BBall   114.  
## 3 Field   111.  
## 4 Field   108.  
## 5 Field   103.  
## 6 WPolo    101  
## 7 BBall    100.
```

# Create new variables from old ones

```
athletes %>%  
  mutate(wt_lb = Wt * 2.2) %>%  
  select(Sport, Sex, Wt, wt_lb) %>%  
  arrange(Wt)
```

```
## # A tibble: 202 x 4  
##   Sport      Sex      Wt wt_lb  
##   <chr>    <chr> <dbl> <dbl>  
## 1 Gym      female  37.8  83.2  
## 2 Gym      female  43.8  96.4  
## 3 Gym      female  45.1  99.2  
## 4 Tennis   female  45.8  101.  
## 5 Tennis   female  47.4  104.  
## 6 Gym      female  47.8  105.  
## 7 T400m    female  49.2  108.  
## 8 Row      female  49.8  110.  
## 9 T400m    female  50.0  110.
```

# Turning the result into a number

Output is always data frame unless you explicitly turn it into something else, eg. the weight of the heaviest athlete, as a number:

```
athletes %>% arrange(desc(Wt)) %>% pluck("Wt", 1)
```

```
## [1] 123.2
```

Or the 20 heaviest weights in descending order:

```
athletes %>%  
  arrange(desc(Wt)) %>%  
  slice(1:20) %>%  
  pluck("Wt")
```

```
## [1] 123.20 113.70 111.30 108.20 102.70 101.00  
## [7] 100.20 98.00 97.90 97.90 97.00 96.90  
## [13] 96.30 94.80 94.80 94.70 94.70 94.60  
## [19] 94.25 94.20
```

## Another way to do the last one

```
athletes %>%  
  arrange(desc(Wt)) %>%  
  slice(1:20) %>%  
  pull("Wt")
```

```
## [1] 123.20 113.70 111.30 108.20 102.70 101.00  
## [7] 100.20 98.00 97.90 97.90 97.00 96.90  
## [13] 96.30 94.80 94.80 94.70 94.70 94.60  
## [19] 94.25 94.20
```

`pull` grabs the column you name *as a vector* (of whatever it contains).

# To find the mean height of the women athletes

Two ways:

```
athletes %>% group_by(Sex) %>% summarize(m = mean(Ht))
```

```
## # A tibble: 2 x 2
##   Sex      m
##   <chr>  <dbl>
## 1 female  175.
## 2 male   186.
```

```
athletes %>%
  filter(Sex == "female") %>%
  summarize(m = mean(Ht))
```

```
## # A tibble: 1 x 1
##       m
##   <dbl>
## 1  175.
```

# Summary of data selection/arrangement “verbs”

Verb	Purpose
<code>select</code>	Choose columns
<code>print</code>	Display non-default # of rows/columns
<code>slice</code>	Choose rows by number
<code>sample_n</code>	Choose random rows
<code>filter</code>	Choose rows satisfying conditions
<code>arrange</code>	Sort in order by column(s)
<code>mutate</code>	Create new variables
<code>group_by</code>	Create groups to summarize by
<code>summarize</code>	Calculate summary statistics (by groups if defined)
<code>pluck</code>	Extract items from data frame
<code>pull</code>	Extract a single column from a data frame as a vector

# Looking things up in another data frame

Recall the tuberculosis data set, tidied:

```
tb3
```

```
## # A tibble: 35,750 x 5
```

```
##   iso2   year gender age   freq
```

```
##   <chr> <dbl> <chr>  <chr> <dbl>
```

```
## 1 AD      1996 m      014     0
```

```
## 2 AD      1996 m     1524     0
```

```
## 3 AD      1996 m     2534     0
```

```
## 4 AD      1996 m     3544     4
```

```
## 5 AD      1996 m     4554     1
```

```
## 6 AD      1996 m     5564     0
```

```
## 7 AD      1996 m      65     0
```

```
## 8 AD      1996 f     014     0
```

```
## 9 AD      1996 f     1524     1
```

```
## 10 AD     1996 f     2534     1
```

# Actual country names

Found actual country names to go with those abbreviations, in spreadsheet:

```
my_url <-  
  "http://www.utsc.utoronto.ca/~butler/c32/ISOCountryCodes081507.xlsx"
```

Note trick for reading in .xlsx from URL:

```
f <- tempfile()  
download.file(my_url, f)  
country_names <- read_excel(f)
```

- set up temporary file
- download spreadsheet to there
- read it from temporary file (which is “local”)



# The country names

```
country_names
```

```
## # A tibble: 252 x 3
##   Code Code_UC Country
##   <chr> <chr>   <chr>
## 1 ad    AD      Andorra
## 2 ae    AE      United Arab Emirates
## 3 af    AF      Afghanistan
## 4 ag    AG      Antigua and Barbuda
## 5 ai    AI      Anguilla
## 6 al    AL      Albania
## 7 am    AM      Armenia
## 8 an    AN      Netherlands Antilles
## 9 ao    AO      Angola
## 10 aq   AQ      Antarctica
## # ... with 242 more rows
```

# Looking up country codes

Matching a variable in one data frame to one in another is called a **join** (database terminology):

```
tb3 %>% left_join(country_names, by = c("iso2" = "Code_UC"))
```

```
## # A tibble: 35,750 x 7
```

##	iso2	year	gender	age	freq	Code	Country
##	<chr>	<dbl>	<chr>	<chr>	<dbl>	<chr>	<chr>
##	1 AD	1996	m	014	0	ad	Andorra
##	2 AD	1996	m	1524	0	ad	Andorra
##	3 AD	1996	m	2534	0	ad	Andorra
##	4 AD	1996	m	3544	4	ad	Andorra
##	5 AD	1996	m	4554	1	ad	Andorra
##	6 AD	1996	m	5564	0	ad	Andorra
##	7 AD	1996	m	65	0	ad	Andorra
##	8 AD	1996	f	014	0	ad	Andorra
##	9 AD	1996	f	1524	1	ad	Andorra

# Total cases by country

```
options(dplyr.summarise.inform=FALSE)
```

```
tb3 %>%  
  group_by(iso2) %>%  
  summarize(cases = sum(freq)) %>%  
  left_join(country_names, by = c("iso2" = "Code_UC")) %>%  
  select(Country, cases)
```

```
## # A tibble: 213 x 2
```

##	Country	cases
##	<chr>	<dbl>
##	1 Andorra	64
##	2 United Arab Emirates	487
##	3 Afghanistan	80005
##	4 Antigua and Barbuda	21
##	5 Anguilla	1
##	6 Albania	2467

or even sorted in order

```
tb3 %>%  
  group_by(iso2) %>%  
  summarize(cases = sum(freq)) %>%  
  left_join(country_names, by = c("iso2" = "Code_UC")) %>%  
  select(Country, cases) %>%  
  arrange(desc(cases))
```

```
## # A tibble: 213 x 2  
##   Country      cases  
##   <chr>      <dbl>  
## 1 China      4065174  
## 2 India      3966169  
## 3 Indonesia  1129015  
## 4 South Africa 900349  
## 5 Bangladesh  758008  
## 6 Vietnam    709695  
## 7 USA        688885
```

# Comments

- This is probably not quite right because of:
  - the 1994-1995 thing
  - there is at least one country in tb3 that was not in country\_names (the NA above). Which?

```
tb3 %>%  
  anti_join(country_names, by = c("iso2" = "Code_UC")) %>%  
  distinct(iso2)
```

```
## # A tibble: 6 x 1  
##   iso2  
##   <chr>  
## 1 CD  
## 2 ME  
## 3 <NA>  
## 4 PS  
## 5 RS  
## 6 TL
```

# Doing things one row at a time

A data frame d:

```
## # A tibble: 3 x 2
##       x1      x2
##   <dbl> <dbl>
## 1     10     13
## 2     11      8
## 3      3      4
```

Want largest value in each row.

# Try number 1

```
d %>% mutate(mx = max(x1, x2))
```

```
## # A tibble: 3 x 3
##       x1      x2      mx
##   <dbl> <dbl> <dbl>
## 1     10     13     13
## 2     11      8     13
## 3      3      4     13
```

- Fails because `max` finds the largest of *all* values in `x1` and `x2` (and repeats answer 3 times), rather than `max` of the two values in each row.

## Try number 2

```
d %>% rowwise() %>%  
  mutate(mx = max(x1, x2))
```

```
## # A tibble: 3 x 3  
##       x1     x2     mx  
##   <dbl> <dbl> <dbl>  
## 1     10     13     13  
## 2     11      8     11  
## 3      3      4      4
```

- Works because rowwise works one row at a time: “find max of all numbers in 1st row”, then “all in 2nd” etc.



# Calculations for groups

- Back to Australian athletes data: suppose we want the correlation between height and weight for athletes of each sport separately:

```
athletes %>% group_by(Sport) %>%  
  summarize(correl = cor(Ht, Wt))
```

```
## # A tibble: 10 x 2  
##   Sport    correl  
##   <chr>    <dbl>  
## 1 BBall    0.880  
## 2 Field    0.471  
## 3 Gym      0.653  
## 4 Netball  0.363  
## 5 Row      0.873  
## 6 Swim     0.860  
## 7 T400m    0.828  
## 8 Tennis   0.826  
## 9 TSprnt   0.806  
## 10 WPolo    0.595
```

## Another way

- Break the data set into groups first:

```
athletes %>% nest_by(Sport)
```

```
## # A tibble: 10 x 2
##   Sport      data
##   <chr>    <list<tibble>>
## 1 BBall    [25 x 12]
## 2 Field    [19 x 12]
## 3 Gym      [4 x 12]
## 4 Netball  [23 x 12]
## 5 Row      [37 x 12]
## 6 Swim     [22 x 12]
## 7 T400m    [29 x 12]
## 8 Tennis   [11 x 12]
## 9 TSprnt   [15 x 12]
## 10 WPolo    [17 x 12]
```

# Comments

- It looks as if all the other variables have disappeared, but they are all hiding in the column data.
- each thing in data is a *dataframe* (inside a dataframe!)
- when a column consists of things that are *not* single numbers, pieces of text, etc., it's called a **list-column** (see `list` in column header).
- to use this, we work `rowwise` on each sport and the data that belongs to that sport.

# The rowwise way

```
athletes %>% nest_by(Sport) %>%  
  mutate(correl = with(data, cor(Ht, Wt)))
```

```
## # A tibble: 10 x 3  
##   Sport      data correl  
##   <chr>    <list<tibble>> <dbl>  
## 1 BBall    [25 x 12]  0.880  
## 2 Field    [19 x 12]  0.471  
## 3 Gym      [4 x 12]   0.653  
## 4 Netball  [23 x 12]  0.363  
## 5 Row      [37 x 12]  0.873  
## 6 Swim     [22 x 12]  0.860  
## 7 T400m    [29 x 12]  0.828  
## 8 Tennis   [11 x 12]  0.826  
## 9 TSprnt   [15 x 12]  0.806  
## 10 WPolo   [17 x 12]  0.595
```

# Another use of list-columns

- Let's find the five-number summary of heights for male and female athletes:

```
athletes %>%  
  group_by(Sex) %>%  
  summarize(q = quantile(Ht))
```

```
## # A tibble: 10 x 2  
##   Sex      q  
##   <chr> <dbl>  
## 1 female 149.  
## 2 female 171.  
## 3 female 175  
## 4 female 180.  
## 5 female 196.  
## 6 male   165.  
## 7 male   180.  
## 8 male   186.  
## 9 male   191  
## 10 male  209.
```

# Better

```
athletes %>% group_by(Sex) %>%  
  summarize(q = list(quantile(Ht)))
```

```
## # A tibble: 2 x 2  
##   Sex      q  
##   <chr> <list>  
## 1 female <dbl [5]>  
## 2 male   <dbl [5]>
```

- each five-number summary is in a list-column, labelled by which Sex it belongs to.

# To get this out

- to look at it:

```
athletes %>% group_by(Sex) %>%  
  summarize(q = list(quantile(Ht))) %>%  
  unnest(q)
```

```
## # A tibble: 10 x 2  
##   Sex      q  
##   <chr> <dbl>  
## 1 female 149.  
## 2 female 171.  
## 3 female 175  
## 4 female 180.  
## 5 female 196.  
## 6 male   165.  
## 7 male   180.  
## 8 male   186.  
## 9 male   191  
## 10 male  209.
```

## Even better

- `quantile` produces a “named vector” with the quantiles labelled:

```
quantile(athletes$Ht)
```

```
##           0%          25%          50%          75%          100%  
## 148.900 174.000 179.700 186.175 209.400
```

- which we turn into dataframe thus:

```
enframe(quantile(athletes$Ht))
```

```
## # A tibble: 5 x 2  
##   name  value  
##   <chr> <dbl>  
## 1 0%      149.  
## 2 25%     174  
## 3 50%     180.  
## 4 75%     186.  
## 5 100%    209.
```



## for males and female athletes

```
athletes %>% group_by(Sex) %>%  
  summarize(q = list(enframe(quantile(Ht)))) %>%  
  unnest(q)
```

```
## # A tibble: 10 x 3  
##   Sex      name value  
##   <chr> <chr> <dbl>  
## 1 female 0%      149.  
## 2 female 25%     171.  
## 3 female 50%     175  
## 4 female 75%     180.  
## 5 female 100%    196.  
## 6 male   0%      165.  
## 7 male   25%     180.  
## 8 male   50%     186.  
## 9 male   75%     191  
## 10 male  100%    209.
```

# Showing off

- we see pivot\_wider later:

```
athletes %>% group_by(Sex) %>%  
  summarize(q = list(enframe(quantile(Ht)))) %>%  
  unnest(q) %>%  
  pivot_wider(names_from = name, values_from = value)
```

```
## # A tibble: 2 x 6  
##   Sex      `0%` `25%` `50%` `75%` `100%`  
##   <chr> <dbl> <dbl> <dbl> <dbl> <dbl>  
## 1 female 149.  171.  175  180.  196.  
## 2 male  165.  180.  186.  191  209.
```

# Simulation

- if I take a sample of 16 observations from a normal distribution with mean 100 and SD 15, what will its mean look like?
- one sample:

```
rnorm(16, 100, 15)
```

```
## [1] 71.42762 110.90188 121.94001 90.40792  
## [5] 85.75109 72.39900 119.84313 108.99881  
## [9] 99.08615 114.11937 81.80634 99.93786  
## [13] 90.21122 117.36629 108.88385 95.17335
```

- do it lots of times (say 1000).
- set up dataframe with 1000 samples numbered, and do it rowwise.


# The simulation

```
tibble(sim = 1:1000) %>%  
  rowwise() %>%  
  mutate(sample = list(rnorm(16, 100, 15))) %>%  
  mutate(sample_mean = mean(sample)) -> d  
d
```

```
## # A tibble: 1,000 x 3  
##       sim sample      sample_mean  
##   <int> <list>         <dbl>  
## 1     1 1 <dbl [16]>         95.3  
## 2     2 2 <dbl [16]>        105.  
## 3     3 3 <dbl [16]>        101.  
## 4     4 4 <dbl [16]>        104.  
## 5     5 5 <dbl [16]>        104.  
## 6     6 6 <dbl [16]>         95.4  
## 7     7 7 <dbl [16]>        102.  
## 8     8 8 <dbl [16]>         97.7
```

# A histogram of the sample means

```
ggplot(d, aes(x = sample_mean)) + geom_histogram(bins = 10)
```



- sample means are closer to 100 than individual observations are
- distribution of sample means is also normal